



Dent.py: Training Dental Models from Zero to Hero

Part 2 of 4

# Formulating the Problem

# Objectives

- Understanding the Problem and Tasks
  - Understand the needs and requirements of problem formulation
  - Explore different deep learning tasks, including classification, semantic segmentation, and instance segmentation.
- How Models Learn
  - Dive into how models learn by understanding losses and optimizers.
  - Evaluate model performance through the introduction of key metrics.




What is the problem?

# Problem Formulation

- How to define your problem?
  - Define the Objective
  - Understand the Input Data
  - Establish metrics for success
  - Select the Learning Paradigm
  - Identify Constraints
  - Set Evaluation Criteria



# Define the Objective



## What is the task?

Classification  
Segmentation



## What are the desired outputs?



Predict lesion presence from X-Ray scans  
Calculate volume of segmentation masks  
from CBCT volumes

# Recap: Tasks



# Classification

- Involve categorizing medical images or specific regions within them into predefined classes or labels
- Examples
  - Caries Detection
  - Lesion Classification
- Task Definition
  - Single Output → Yes or No
    - Does this image contain caries?
  - Multiple Outputs → Mild, Moderate, Severe
    - What is the level of severity of caries in these images?



# Segmentation

- Involve partitioning an image into distinct regions or structures
- Examples:
  - Hard Tissue Segmentation
  - Pulp Instance Segmentation
- Task Definition
  - Single Output → Matches input size (Semantic)
    - Detect all lesions in the mouth
  - Multiple Outputs → Matches input size (Instance)
    - Detect all lesions in the mouth with their types respectively





# Understand the Input Data



## What data is available?

Images (2D X-rays, 3D CT, Intraoral RGB)  
Text (reports, diagnosis)



## Is the data labeled or unlabeled?

3D CBCT scans with annotated segmentations.  
2D intraoral images with bounding boxes

# Success Metrics



## Map Inputs to Outputs

Define the relationship between input data and output label

Example

- Input: CBCT scan → Output: Semantic segmentation of pulp and IAN.



## Choose Appropriate Metrics


Measures Input-Output mapping success

Example

- Accuracy of segmentation of pulp
- Precision of lesion classification


# Choosing the Right Metric

- Segmentation Tasks
  - Dice
  - IoU
  - HD95




# Choosing the Right Metric

- Classification Metrics
  - Precision
  - Accuracy
  - Recall



# Choosing the Right Metric

- Honorable Mentions
  - MAE
  - Specificity / Sensitivity



# Learning Paradigm


Supervised: Requires labeled data.

Semi-Supervised: Requires only a subset of labeled data.

Unsupervised: Explores hidden structures in data.

Reinforcement: Learns through rewards and penalties.

# Identify Constraints



## Compute resources

GPU/CPU availability  
Memory and Disk requirements



## Real-world limitations

Label scarcity  
Noisy data  
Class Imbalance

# Set Evaluation Criteria



**Use benchmarks to assess progress.**



**Define metrics and comparison baselines.**

Example: Compare segmentation results with state-of-the-art models.



**Visual Representation**

A diagram showcasing how raw data transitions into a model output can help clarify the formulation.



# How Models Learn



# Key Concept

---

## What is a model

- A nonlinear function transforming some input into output
- Our goal is trying to fit the **model weights**

$$f(x; \theta) = y$$

*$f$  : The model (nonlinear function)*

*$x$  : The input data*

*$\theta$  : The parameters of the model (e. g., weights and biases)*

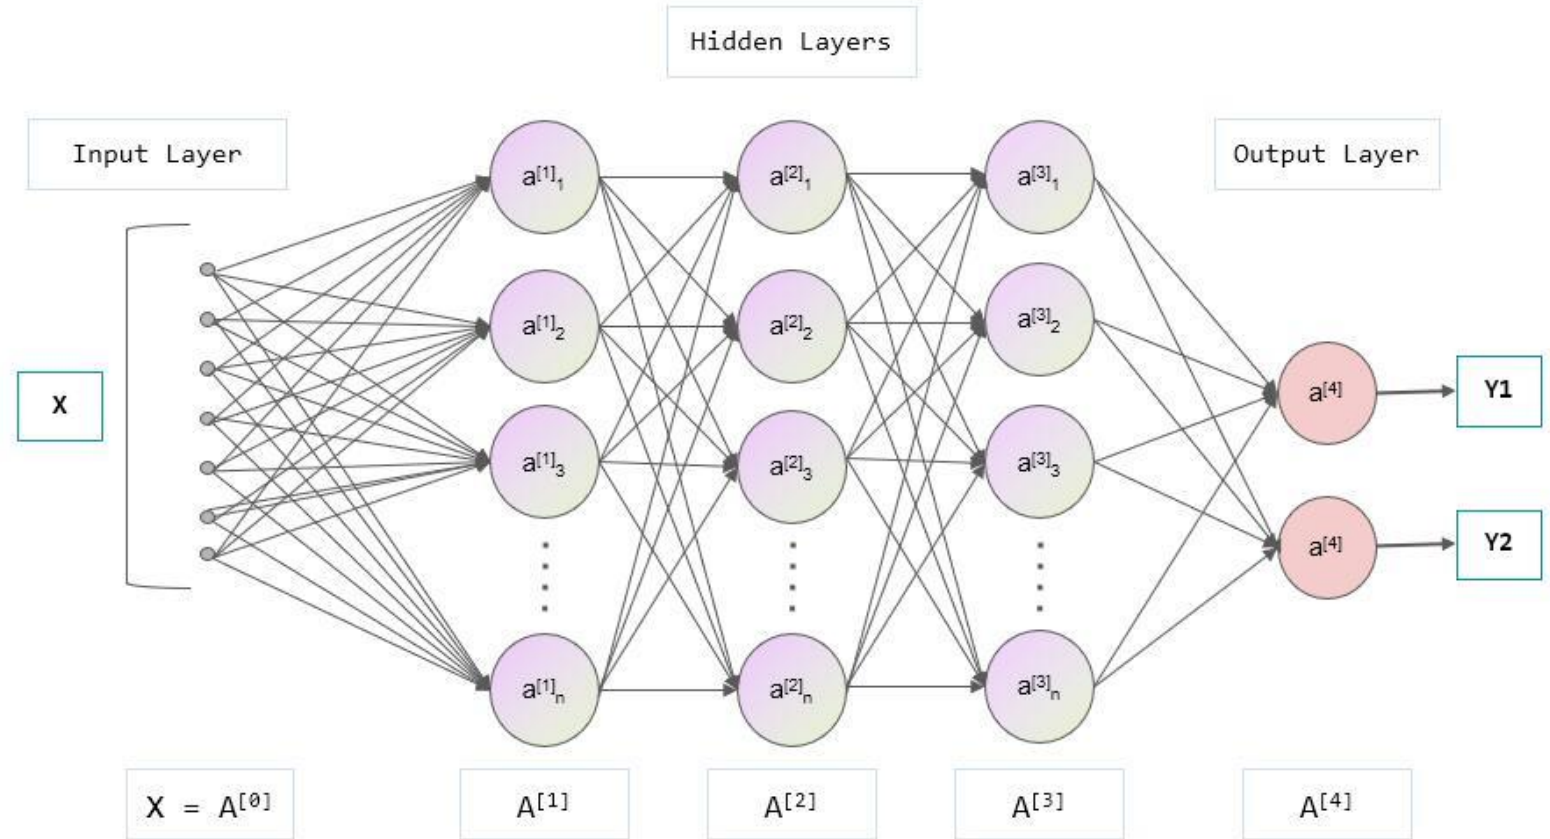
*$y$  : The output (prediction).*

# Training Process

- Goal
  - Adjust model parameters (weights and biases) to minimize error between predictions and true outputs.
- Core Steps:
  - Forward Pass
    - Calculate predictions.
  - Loss Calculation
    - Measure the difference between predictions and true labels using a loss function (e.g., MSE, Cross-Entropy).
  - Backward Pass
    - Compute gradients to update parameters using backpropagation.

# Simple Network

$$Y = F(x) \rightarrow WX + b$$



# Backpropagation

---


- Definition
  - A method for computing gradients of the loss function with respect to model parameters by applying the chain rule.
- Key Concepts:
  - Gradients flow backward from the output layer to earlier layers.
  - Updates are proportional to the gradient and learning rate.

$$w \leftarrow w - \eta \cdot \frac{\partial w}{\partial L}$$

*where  $\eta$  is the learning rate,  $L$  is the loss, and  $w$  are the weights.*

# Optimizers


Algorithms to update weights efficiently based on gradients



Optimizer	Key Features
<b>SGD</b>	Simple, uses the full gradient. Can be slow for large datasets.
<b>Momentum</b>	Adds velocity to updates, reducing oscillations.
<b>RMSProp</b>	Scales learning rate based on recent gradients.
<b>Adam</b>	Combines Momentum and RMSProp, adapts learning rates for each parameter.

# Learning Rate

- A critical hyperparameter that controls the step size in updates.
  - Too high: May overshoot minima.
  - Too low: Converges slowly or gets stuck.



# Loss Functions

Measure how wrong the model's predictions are compared to the actual answers

Task Friendly

Metrics Compatible



Example

Using a **CrossEntropy** loss for **Segmentation Task** while measuring **Dice Score**



# Loss Functions

CrossEntropy

MSE

L1

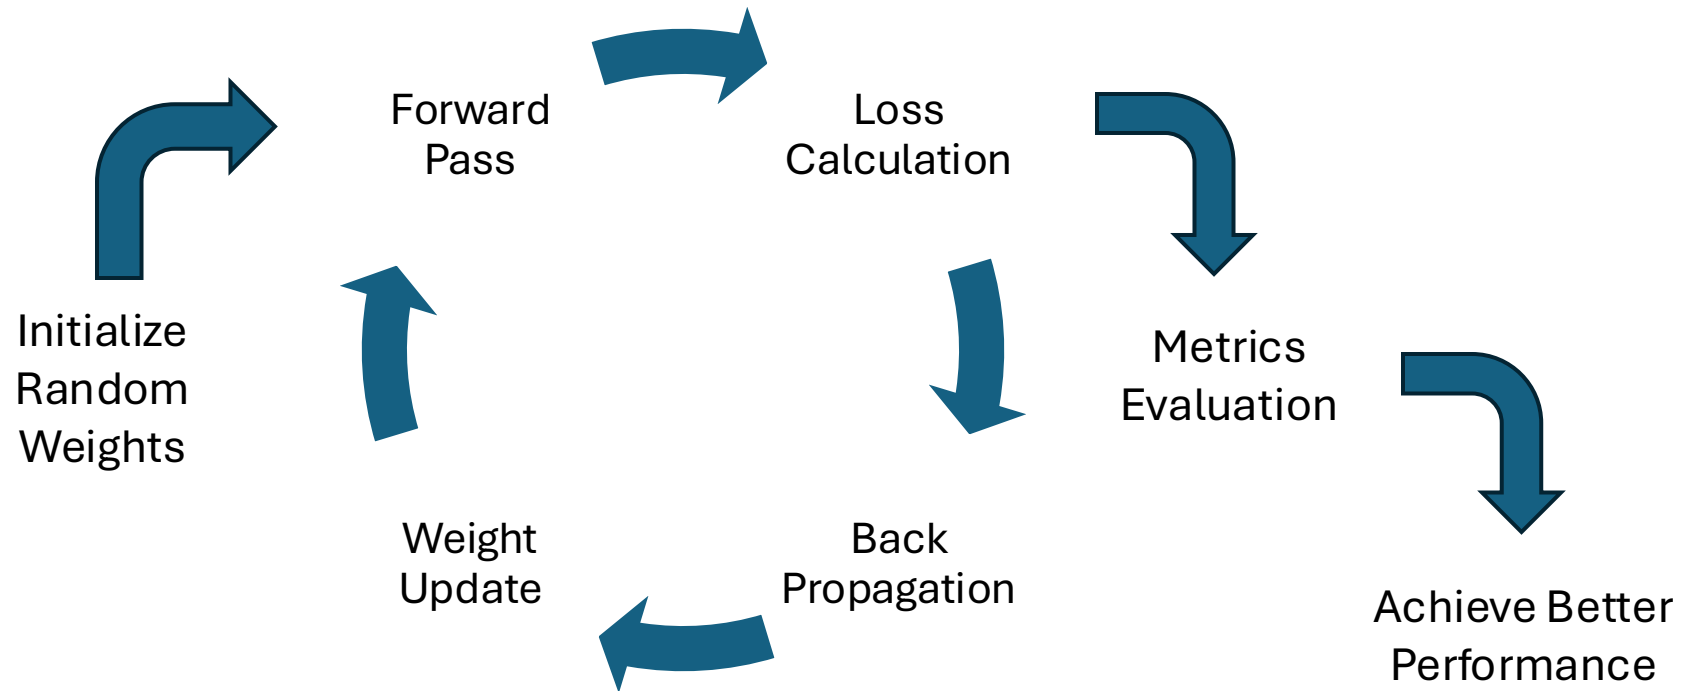
Focal Loss

Combined Losses

- DiceCrossEntropy
- Jaccard-Focal Loss


# Intuition for Learning

---



# Summary

---



# Questions?

---

# Workshop Activity

---

- Notebooks Link
  - <https://github.com/KnightsLab/EMRA-Workshop>

