




Dent.py: Training Dental Models from Zero to Hero

Part 4 of 5

YOLO: You Only Look Once – Deep Dive

Object Detection

- What is object detection?
 - Identifying and locating objects in an image.
 - Used in various applications like autonomous driving, security, retail, etc...



YOLO

- YOLO
 - You Only Look Once
- Overview
 - A fast and efficient object detection algorithm.
 - Processes an image in a single pass, unlike traditional methods.
- Why YOLO is important
 - Real-time processing capability.
 - High accuracy and efficiency.






How It Works?



Grid-based prediction

- The image is divided into an $S \times S$ grid (e.g., 7×7 , 19×19).
- Each grid cell predicts bounding boxes, confidence scores, and class probabilities.
- A grid cell is responsible for detecting objects whose center falls within it.




Bounding box prediction

- Each grid cell predicts multiple bounding boxes with (x, y, w, h) coordinates.
- Predefined boxes of different aspect ratios and sizes improve detection.
- Helps in detecting objects of various scales and shapes.
- Each grid cell predicts adjustments to these anchor boxes rather than free-form bounding boxes.




Confidence Scores

- Represents how likely an object is present and the accuracy of the bounding box.
- Confidence score = Object probability \times IoU (Intersection over Union).
- Higher confidence scores indicate more reliable predictions.




Non-Maximum Suppression


- Removes overlapping boxes to keep only the most relevant detections.
- Ensures that the best bounding box for each object is retained.



Step 1: Selecting Bounding box with highest score



Step 3: Delete Bounding box with high overlap



Step 5: Final Output



Architecture




Overview

- Input and Feature Extraction
 - Uses a CNN backbone (e.g., Darknet-53 in YOLOv3) to extract hierarchical features.
 - Processes input images through convolutional layers with increasing depth.
 - Detection Head
 - Splits into multiple scales for multi-scale detection (FPN-like structure in YOLOv3).
 - Outputs bounding box coordinates, objectness score, and class probabilities.
-



Building Blocks

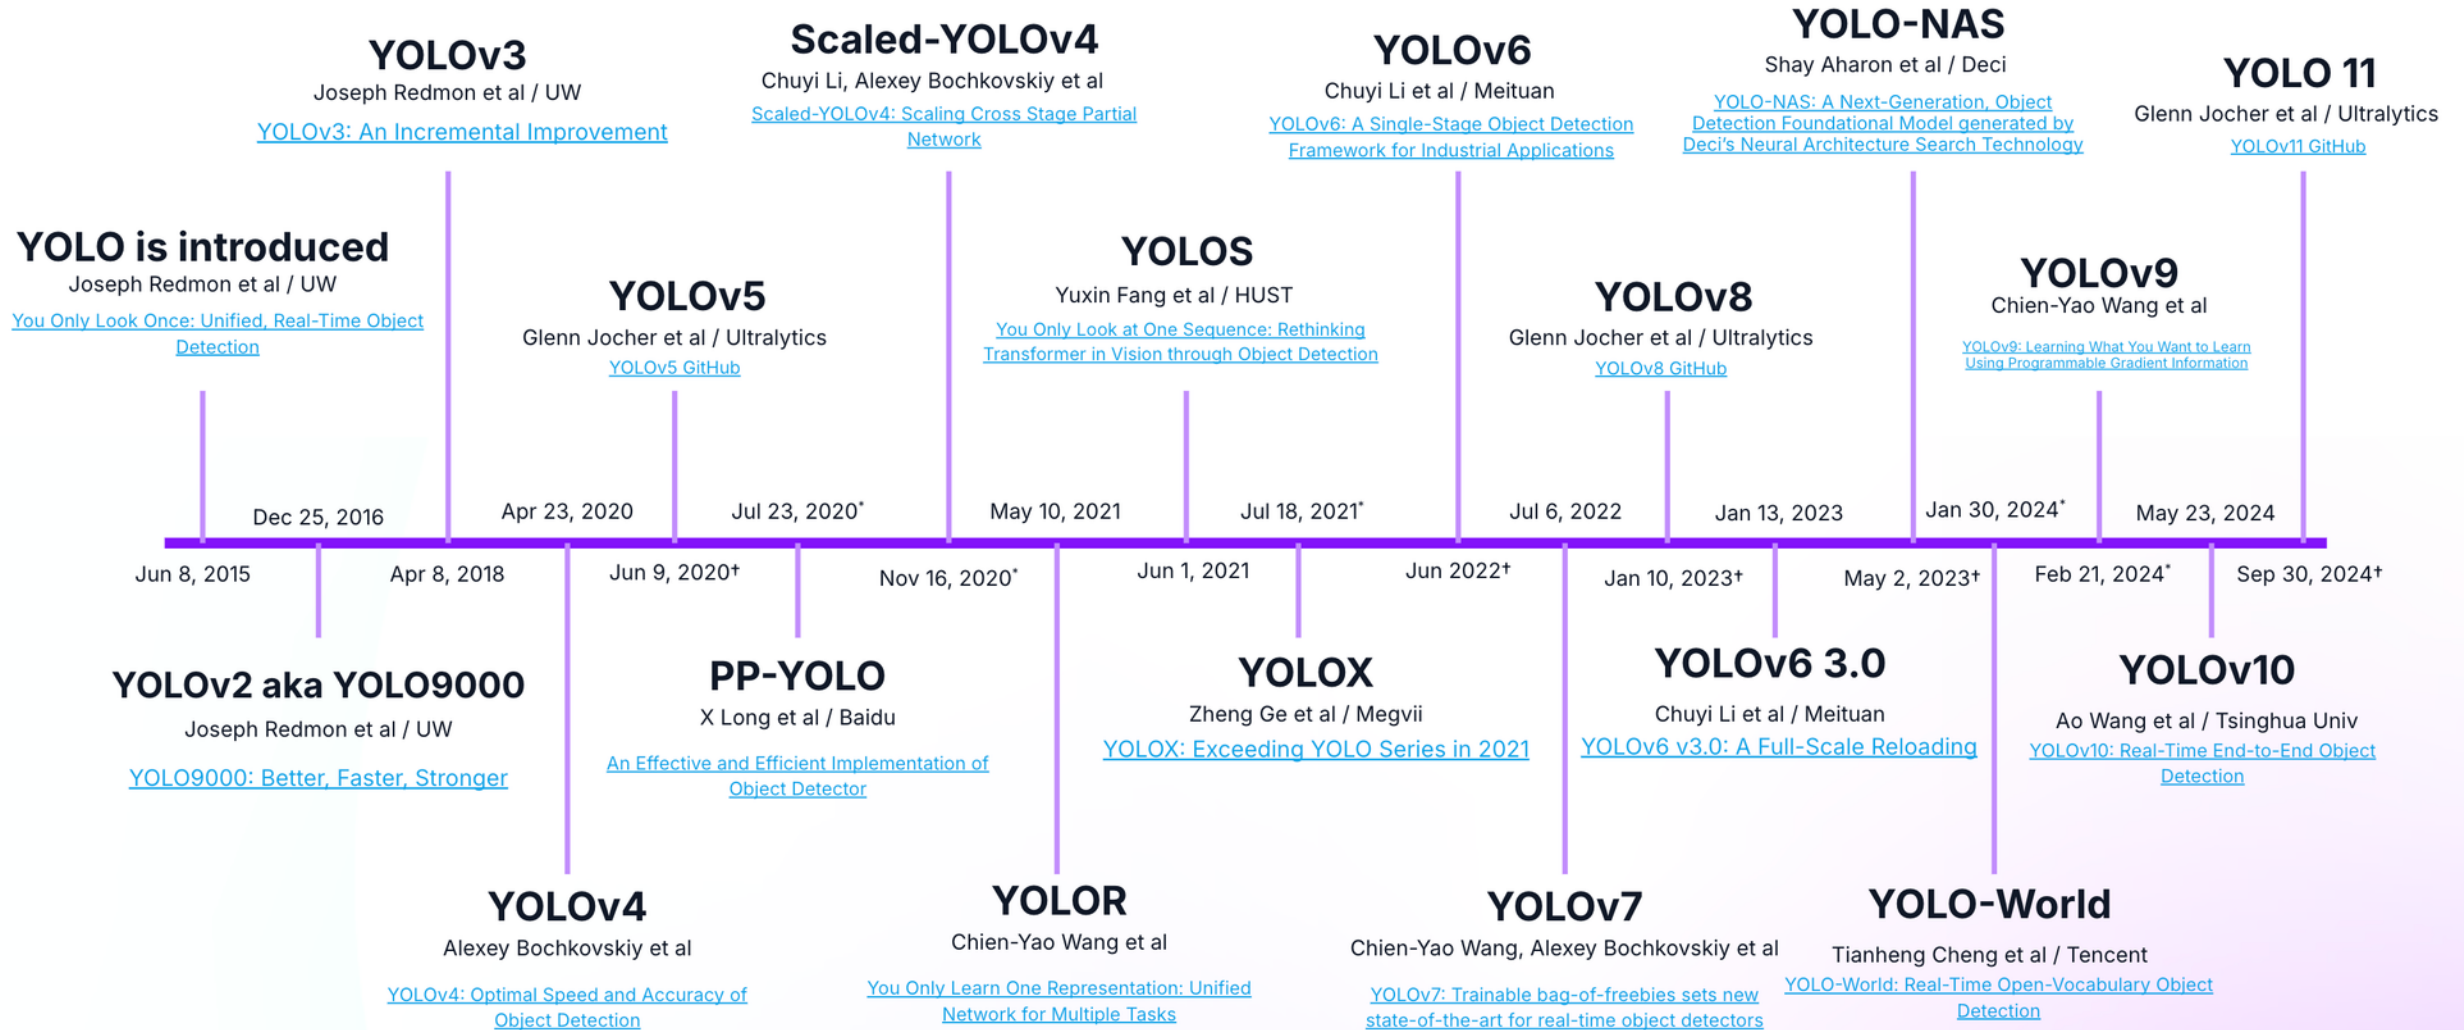
- Convolutional Layers
 - Extract spatial features from images.
 - Batch Normalization
 - Speeds up training and stabilizes learning.
 - Residual Connections
 - Help gradient flow and improve learning efficiency.
 - Leaky ReLU Activation
 - Prevents vanishing gradients and speeds up training.
-







Greatest Hits



* Denotes paper updated after first publication date
* Denotes repository predates paper publication date



YOLOv12

- Released on February 18, 2025
- Introduced in the paper “YOLOv12: Attention-Centric Real-Time Object Detectors.”
- **Key Features**
 - Incorporates advanced attention mechanisms for improved detection accuracy.
 - Optimized for real-time applications with lower latency.
 - Open-source implementation available for fine-tuning and customization.
- **Performance**
 - Benchmarked on the Microsoft COCO dataset.
 - Achieves higher mean Average Precision (mAP) while reducing computational overhead.

The background of the slide is an abstract pattern consisting of a grid of small squares. The squares are colored in shades of red, green, and blue, arranged in a way that creates a sense of depth and perspective, as if the grid is receding into the distance. The colors are slightly blurred, giving the pattern a soft, painterly quality.

Applications

General



Autonomous Vehicles

Detects pedestrians, vehicles, traffic signs in real-time.



Security & Surveillance

Identifies objects in CCTV footage.



Medical Imaging

Used for detecting tumors, anomalies in medical scans.



Retail & Logistics

Automated checkout systems, inventory tracking.

Dentistry



Tooth Segmentation

Identifying individual teeth for orthodontic planning.



Caries and Cavity Detection

Detecting early-stage cavities in dental X-rays.



Root Canal Detection

Assisting endodontists in visualizing canal morphology.



IAN (Inferior Alveolar Nerve) Localization

Preventing nerve damage during surgeries.

The background of the image is a complex 3D maze. The maze is constructed from thick, light-colored walls that create a series of concentric, winding paths. The perspective is from a slightly elevated angle, looking down into the maze. The lighting is soft, casting gentle shadows that emphasize the three-dimensional nature of the structure. The overall color palette is muted, consisting of various shades of gray and beige. Centered within this intricate labyrinthine pattern is the text "Pros and Cons" in a clean, white, sans-serif font. The text is clearly legible against the darker, more complex background of the maze.

Pros and Cons

Pros

- Fast and efficient, capable of real-time detection.
- Single-stage processing, end-to-end learning.
- Works well in real-world applications.

Cons

- Struggles with detecting small objects.
- May not be as accurate as two-stage methods like Faster R-CNN.
- Limited interpretability due to end-to-end learning.

Questions?

Workshop Activity

- Notebooks Link
 - <https://github.com/KnightsLab/EMRA-Workshop>

